

Replication file for the Russian intermarriage paper

Alexey Bessudnov

5 February 2020

The file has been compiled with R version 3.6.2 and R Markdown version 2.0.

```
# Attaching the packages.
library(tidyverse)
# tidyverse version 1.3.0
library(knitr)
# knitr version 1.26
library(gnm)
# gnm version 1.1.0
library(lmtest)
# lmtest version 0.9.37
```

Read the data

The data contain six tables for the following cities: Moscow, Rostov, Kazan, Ufa, Makhachkala and Vladikavkaz. For each city the data represent the number of married couples by wife's and husband's ethnicities, only for locally born wives and for wives in three age groups (16 to 35, 36 to 50 and over 50), taken from the 2010 Russian census.

```
Kazan <- read_csv("data/Kazan.csv")
Moscow <- read_csv("data/Moscow.csv")
Vladikavkaz <- read_csv("data/Vladikavkaz.csv")
Makhachkala <- read_csv("data/Makhachkala.csv")
Ufa <- read_csv("data/Ufa.csv")
Rostov <- read_csv("data/Rostov.csv")
```

Prepare the data

First we want to collapse the data tables combining marriages across all three age groups.

```
# function to collapse data tables
collapseData <- function(df) {
  df %>%
    group_by(ethn.wife, ethn.husband) %>%
    summarise(
      Freq = sum(Freq)
    )
}
```

Visualise contingency tables for intermarriages

The heatmaps represent logged frequencies in the contingency tables of ethnicities of wives and husbands. Grey cells are cells with zero observations (so that log is not defined). Note that the data are for locally born women only.

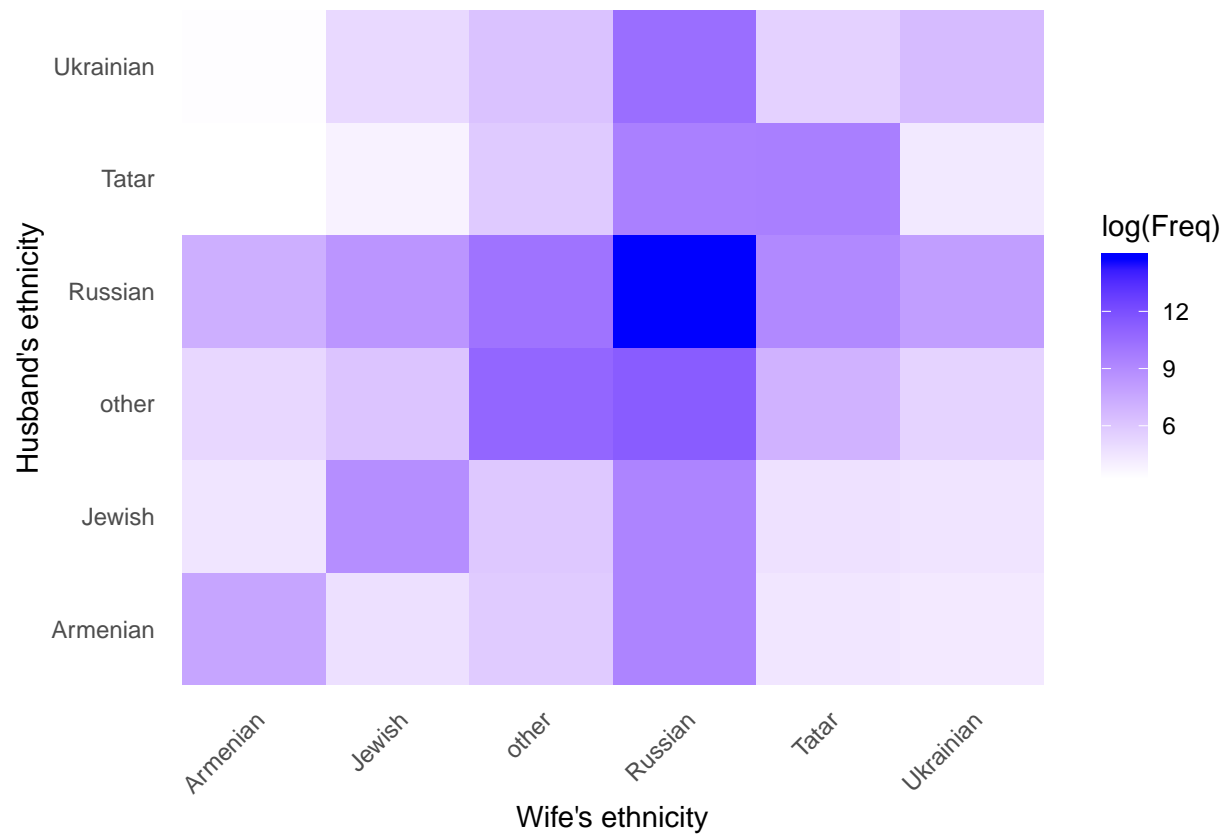
This is Figure 1 in the paper.

```
# function to produce heat maps
heatMap <- function(df){
  collapseData(df) %>%
    ggplot(aes(x = ethn.wife, y = ethn.husband, fill = log(Freq))) +
    geom_tile() +
    xlab("Wife's ethnicity") +
    ylab("Husband's ethnicity") +
    scale_fill_gradient(low = "white", high = "blue") +
    theme_classic() +
    theme(axis.text.x = element_text(angle = 45, hjust = 1),
          axis.line = element_blank(),
          axis.ticks = element_blank())
}


# combine data tables for six cities into a list
sixCities <- list(Moscow, Rostov, Kazan, Ufa, Vladikavkaz, Makhachkala)
# assigning names to the elements
names(sixCities) <- c("Moscow", "Rostov", "Kazan", "Ufa", "Vladikavkaz", "Makhachkala")

sixCities %>% map(heatMap)


## $Moscow
```




\$Rostov




\$Kazan




\$Ufa



\$Vladikavkaz



\$Makhachkala



Percent by ethnic group for men and women

See Table 1 in the paper.

Distribution of cases by ethnic group, separately for women and men.

Functions to summarise percentages by ethnic group for women and men

```
percentByGroupWife <- function(df){
  collapseData(df) %>%
    group_by(ethn.wife) %>%
    summarise(
      n = sum(Freq)
    ) %>%
    mutate(perc = n / sum(n) * 100)
}

percentByGroupHusband <- function(df){
  collapseData(df) %>%
    group_by(ethn.husband) %>%
    summarise(
      n = sum(Freq)
    ) %>%
    mutate(perc = n / sum(n) * 100)
}
```


Percentages for women.

```
sixCities %>% map(percentByGroupWife)
```

```
## $Moscow
## # A tibble: 6 x 3
##   ethn.wife      n    perc
##   <chr>         <dbl> <dbl>
## 1 Armenian     3930  0.140
## 2 Jewish      12656  0.452
## 3 other        79646  2.85
## 4 Russian    2673121 95.5
## 5 Tatar        25738  0.919
## 6 Ukrainian    4350  0.155
##
## $Rostov
## # A tibble: 4 x 3
##   ethn.wife      n    perc
##   <chr>         <dbl> <dbl>
## 1 Armenian     7156  2.44
## 2 other        6389  2.18
## 3 Russian    278585 94.9
## 4 Tatar        1282  0.437
##
## $Kazan
## # A tibble: 4 x 3
##   ethn.wife      n    perc
##   <chr>         <dbl> <dbl>
## 1 Chuvash     1029  0.313
## 2 other       2275  0.692
## 3 Russian   174371 53.0
## 4 Tatar     151214 46.0
##
## $Ufa
## # A tibble: 7 x 3
##   ethn.wife      n    perc
##   <chr>         <dbl> <dbl>
## 1 Bashkir    39342 13.1
## 2 Chuvash    2122  0.707
## 3 Mari       2287  0.762
## 4 other      5667  1.89
## 5 Russian   170522 56.8
## 6 Tatar      77160 25.7
## 7 Ukrainian  2862  0.954
##
## $Vladikavkaz
## # A tibble: 5 x 3
##   ethn.wife      n    perc
##   <chr>         <dbl> <dbl>
## 1 Armenian    2829  2.82
## 2 Georgian    1841  1.83
## 3 Ossetian   66814 66.5
## 4 other       2716  2.70
## 5 Russian   26244 26.1
##
```

```
## $Makhachkala
## # A tibble: 13 x 3
##   ethn.wife      n    perc
##   <chr>         <dbl> <dbl>
## 1 Aghul         2936  0.875
## 2 Avar          79603 23.7
## 3 Azerbaijani  2354  0.702
## 4 Dargin        49991 14.9
## 5 Kumyk         74005 22.1
## 6 Laki          42517 12.7
## 7 Lezgian       44795 13.4
## 8 Nogai         4129  1.23
## 9 other         4119  1.23
## 10 Russian      20164  6.01
## 11 Rutulian     3689  1.10
## 12 Tabasaran    6051  1.80
## 13 Tsakhur      1149  0.342
```

Percentages for men.

```
sixCities %>% map(percentByGroupHusband)
```

```
## $Moscow
## # A tibble: 6 x 3
##   ethn.husband      n    perc
##   <chr>         <dbl> <dbl>
## 1 Armenian      14709  0.525
## 2 Jewish        19143  0.684
## 3 other         139628  4.99
## 4 Russian       2559553 91.4
## 5 Tatar         29708  1.06
## 6 Ukrainian     36700  1.31
##
```

```
## $Rostov
## # A tibble: 4 x 3
##   ethn.husband      n    perc
##   <chr>         <dbl> <dbl>
## 1 Armenian      10626  3.62
## 2 other         15313  5.22
## 3 Russian       265987 90.7
## 4 Tatar         1486  0.506
##
```

```
## $Kazan
## # A tibble: 4 x 3
##   ethn.husband      n    perc
##   <chr>         <dbl> <dbl>
## 1 Chuvash       2071  0.630
## 2 other         8349  2.54
## 3 Russian      170172 51.7
## 4 Tatar        148297 45.1
##
```

```
## $Ufa
## # A tibble: 7 x 3
##   ethn.husband      n    perc
##   <chr>         <dbl> <dbl>
```

```
## 1 Bashkir      39504 13.2
## 2 Chuvash      2630  0.877
## 3 Mari         2512  0.837
## 4 other        11226 3.74
## 5 Russian     163826 54.6
## 6 Tatar        75654 25.2
## 7 Ukrainian    4610  1.54
##
## $Vladikavkaz
## # A tibble: 5 x 3
##   ethn.husband    n perc
##   <chr>          <dbl> <dbl>
## 1 Armenian      3326  3.31
## 2 Georgian      2132  2.12
## 3 Ossetian      67945 67.6
## 4 other          4045  4.03
## 5 Russian      22996 22.9
##
## $Makhachkala
## # A tibble: 13 x 3
##   ethn.husband    n perc
##   <chr>          <dbl> <dbl>
## 1 Aghul          3058  0.911
## 2 Avar           83075 24.8
## 3 Azerbaijani    2795  0.833
## 4 Dargin         50893 15.2
## 5 Kumyk          73408 21.9
## 6 Laki           42424 12.6
## 7 Lezgian        44316 13.2
## 8 Nogai          3944  1.18
## 9 other          4367  1.30
## 10 Russian       15718  4.68
## 11 Rutulian       3924  1.17
## 12 Tabasaran      6299  1.88
## 13 Tsakhur       1281  0.382
```

Percent married within group

See Table 1 in the paper.

The tables below show the percentages married within group for the ethnic groups in six cities, separately for men and women (locally born women only).

```
# a function to produce tables with % married within group
marriedWithin <- function(df){
  # number of women
  df1 <- df %>%
    group_by(ethn.wife) %>%
    summarise(
      nWomen = sum(Freq)
    )
  # number of men
  df2 <- df %>%
    group_by(ethn.husband) %>%
```

```

      summarise(
        nMen = sum(Freq)
      ) %>%
      rename(ethn.wife = ethn.husband)

# number of people married within their group

df3 <- df %>%
  filter(ethn.wife == ethn.husband) %>%
  group_by(ethn.wife) %>%
  summarise(
    nWithin = sum(Freq)
  )

# put it all together

df1 %>%
  left_join(df2, by = c("ethn.wife")) %>%
  left_join(df3, by = c("ethn.wife")) %>%
  mutate(percWomenWithin = nWithin / nWomen * 100) %>%
  mutate(percMenWithin = nWithin / nMen * 100) %>%
  rename(ethnicity = ethn.wife) %>%
  select(-nWithin) %>%
  select(ethnicity, percWomenWithin, nWomen, percMenWithin, nMen) %>%
  filter(ethnicity != "other")
}

sixCities %>% map(marriedWithin)

## $Moscow
## # A tibble: 5 x 5
##   ethnicity percWomenWithin  nWomen percMenWithin   nMen
##   <chr>          <dbl>   <dbl>         <dbl> <dbl>
## 1 Armenian      57.4    3930          15.3  14709
## 2 Jewish        55.0   12656          36.3   19143
## 3 Russian       94.0 2673121          98.2 2559553
## 4 Tatar         58.4   25738          50.6   29708
## 5 Ukrainian     17.0    4350           2.02   36700
##
## $Rostov
## # A tibble: 3 x 5
##   ethnicity percWomenWithin  nWomen percMenWithin   nMen
##   <chr>          <dbl>   <dbl>         <dbl> <dbl>
## 1 Armenian      66.1    7156          44.5  10626
## 2 Russian       93.4 278585          97.9 265987
## 3 Tatar         46.3   1282          40.0   1486
##
## $Kazan
## # A tibble: 3 x 5
##   ethnicity percWomenWithin  nWomen percMenWithin   nMen
##   <chr>          <dbl>   <dbl>         <dbl> <dbl>
## 1 Chuvash       26.1   1029          13.0   2071
## 2 Russian       81.4 174371          83.4 170172
## 3 Tatar         80.1 151214          81.7 148297

```

```
##
## $Ufa
## # A tibble: 6 x 5
##   ethnicity percWomenWithin nWomen percMenWithin nMen
##   <chr>          <dbl>   <dbl>          <dbl> <dbl>
## 1 Bashkir        52.5   39342          52.2  39504
## 2 Chuvash        24.4    2122          19.7   2630
## 3 Mari           39.8    2287          36.2   2512
## 4 Russian        76.8  170522          79.9  163826
## 5 Tatar          55.9   77160          57.0   75654
## 6 Ukrainian       8.91   2862           5.53   4610
##
## $Vladikavkaz
## # A tibble: 4 x 5
##   ethnicity percWomenWithin nWomen percMenWithin nMen
##   <chr>          <dbl>   <dbl>          <dbl> <dbl>
## 1 Armenian        61.9   2829           52.7   3326
## 2 Georgian        52.6   1841           45.4   2132
## 3 Ossetian        92.7  66814           91.2  67945
## 4 Russian         74.9  26244           85.5  22996
##
## $Makhachkala
## # A tibble: 12 x 5
##   ethnicity percWomenWithin nWomen percMenWithin nMen
##   <chr>          <dbl>   <dbl>          <dbl> <dbl>
## 1 Aghul           86.3   2936           82.9   3058
## 2 Avar            92.4  79603           88.5  83075
## 3 Azerbaijani     72.9   2354           61.4   2795
## 4 Dargin          89.9  49991           88.3  50893
## 5 Kumyk           91.4  74005           92.2  73408
## 6 Laki            88.3  42517           88.5  42424
## 7 Lezgian         88.3  44795           89.2  44316
## 8 Nogai           86.1   4129           90.1   3944
## 9 Russian         71.4  20164           91.6  15718
## 10 Rutulian       80.8   3689           75.9   3924
## 11 Tabasaran      85.4   6051           82.0   6299
## 12 Tsakhur        94.6   1149           84.9   1281
```

Odds ratios for ethnic endogamy across six cities

See Table 1 in the paper.

```
# a function to collapse data tables to 2x2 tables
collapse2x2 <- function(df, ethn){
  df %>%
    mutate(ethn2F = ifelse(ethn.wife == ethn, 1, 0)) %>%
    mutate(ethn2M = ifelse(ethn.husband == ethn, 1, 0)) %>%
    group_by(ethn2F, ethn2M) %>%
    summarise(
      Freq = sum(Freq)
    )
}

# a function to calculate log odds ratios with 95% confidence intervals
```

```

oddsRatio <- function(df){
  # freq is a vector of counts
  freq <- df %>% pull(Freq)
  # a is the number of marriages where both husband and wife are not from group i
  # as.numeric() added to avoid integer overflow
  a <- as.numeric(freq[1])
  # b and c are the number of intermarriages
  b <- as.numeric(freq[2])
  c <- as.numeric(freq[3])
  # d is the number of endogamous intermarriages for group i
  d <- as.numeric(freq[4])
  # calculating log oddsratio
  logOR <- log((a*d) / (b*c))
  # calculating the standard error for log odds ratio (see https://www.ncbi.nlm.nih.gov/pmc/articles/
  se <- sqrt(1/a + 1/b + 1/c + 1/d)
  # calculating 95% confidence intervals
  lowerCI <- logOR - 1.96*se
  upperCI <- logOR + 1.96*se
  return(c(logOR, lowerCI, upperCI))
}

# a function to loop over all ethnic groups and age groups to produce a data frame with log ORs and CIs
ethnOR <- function(df){
  # Initialise a data frame for the results
  n <- length(levels(df$ethn.wife)) * 3 # 3 age groups
  results <- data.frame(city = rep(NA, n),
                        ethnGroup = rep(NA, n),
                        ageGroup = rep(NA, n),
                        logOR = rep(NA, n),
                        lowerCI = rep(NA, n),
                        upperCI = rep(NA, n))

  k <- 1
  for (i in levels(as.factor(df$ethn.wife))) {
    for (j in levels(as.factor(df$age.wife))) {
      dfNew <- filter(df, age.wife == j)
      res <- collapse2x2(dfNew, i) %>%
        oddsRatio()
      results[k,1] <- df$city[1]
      results[k,2] <- i
      results[k,3] <- j
      results[k,4] <- res[1]
      results[k,5] <- res[2]
      results[k,6] <- res[3]
      k <- k + 1
    }
  }
  results
}

# same function as above (to produce ORs), but collapsing all age groups together.
ethnOR2 <- function(df){
  # Initialise a data frame for the results

```

```

n <- length(as.factor(levels(df$ethn.wife)))
results <- data.frame(city = rep(NA, n),
                      ethnGroup = rep(NA, n),
                      logOR = rep(NA, n),
                      lowerCI = rep(NA, n),
                      upperCI = rep(NA, n))

k <- 1
for (i in levels(as.factor(df$ethn.wife))) {
  dfNew <- collapseData(df)
  res <- collapse2x2(dfNew, i) %>%
    oddsRatio()
  results[k,1] <- df$city[1]
  results[k,2] <- i
  results[k,3] <- res[1]
  results[k,4] <- res[2]
  results[k,5] <- res[3]
  k <- k + 1
}
results
}

# produce a table with ORs for all cities

sixCities %>%
  map_df(ethnOR2)

```

##	city	ethnGroup	logOR	lowerCI	upperCI
## 1	Moscow	Armenian	5.705479	5.639857	5.771101
## 2	Moscow	Jewish	5.627833	5.588552	5.667114
## 3	Moscow	other	3.966289	3.950347	3.982232
## 4	Moscow	Russian	3.332329	3.319777	3.344881
## 5	Moscow	Tatar	5.576025	5.546401	5.605649
## 6	Moscow	Ukrainian	2.755451	2.675678	2.835224
## 7	Rostov	Armenian	4.530279	4.474947	4.585610
## 8	Rostov	other	3.045851	2.993544	3.098159
## 9	Rostov	Russian	3.119585	3.083279	3.155890
## 10	Rostov	Tatar	5.641520	5.513571	5.769470
## 11	Kazan	Chuvash	4.159571	4.013013	4.306130
## 12	Kazan	other	3.312038	3.224746	3.399331
## 13	Kazan	Russian	2.977552	2.959879	2.995225
## 14	Kazan	Tatar	3.110666	3.092584	3.128748
## 15	Ufa	Bashkir	2.648479	2.623760	2.673197
## 16	Ufa	Chuvash	3.811524	3.703621	3.919428
## 17	Ufa	Mari	4.805135	4.708067	4.902204
## 18	Ufa	other	2.645712	2.586156	2.705268
## 19	Ufa	Russian	2.271800	2.254981	2.288619
## 20	Ufa	Tatar	2.001979	1.983535	2.020423
## 21	Ufa	Ukrainian	1.883285	1.751248	2.015323
## 22	Vladikavkaz	Armenian	4.597734	4.506958	4.688510
## 23	Vladikavkaz	Georgian	4.530661	4.422453	4.638868
## 24	Vladikavkaz	Ossetian	4.079002	4.038560	4.119445
## 25	Vladikavkaz	other	3.690593	3.605869	3.775318
## 26	Vladikavkaz	Russian	4.148901	4.104361	4.193441
## 27	Makhachkala	Aghul	8.292627	8.156925	8.428330

```
## 28 Makhachkala      Avar  5.747753  5.714519  5.780987
## 29 Makhachkala Azerbaijani 6.721803  6.612975  6.830630
## 30 Makhachkala      Dargin 6.028406  5.989666  6.067145
## 31 Makhachkala      Kumyk  6.166029  6.129315  6.202743
## 32 Makhachkala      Laki   6.098798  6.057875  6.139722
## 33 Makhachkala      Lezgian 6.112423  6.071834  6.153013
## 34 Makhachkala      Nogai  8.565135  8.432380  8.697889
## 35 Makhachkala      other  5.525981  5.449300  5.602662
## 36 Makhachkala      Russian 6.386929  6.324836  6.449023
## 37 Makhachkala      Rutulian 7.292368  7.188555  7.396180
## 38 Makhachkala      Tabasaran 7.435590  7.343424  7.527756
## 39 Makhachkala      Tsakhur 10.315557 10.023480 10.607634
```

```
# produce a data frame with log ORs for all six cities; remove missing values;
# reorder levels for age groups
```

```
ORData <- sixCities %>%
  map_df(ethnOR) %>%
  filter(is.finite(logOR)) %>%
  mutate(ageGroup = fct_relevel(ageGroup, "16-35", "36-50", ">50"))
```

```
head(ORData)
```

```
##      city ethnGroup ageGroup   logOR lowerCI upperCI
## 1 Moscow  Armenian    >50 5.152011 5.018191 5.285831
## 2 Moscow  Armenian   16-35 6.014574 5.906776 6.122372
## 3 Moscow  Armenian   36-50 5.740118 5.630176 5.850060
## 4 Moscow   Jewish    >50 5.212030 5.159099 5.264961
## 5 Moscow   Jewish   16-35 6.073428 5.976779 6.170078
## 6 Moscow   Jewish   36-50 5.702149 5.626474 5.777825
```


```
# a function to produce dot plots with confidence intervals
```

```
plotOR <- function(x){
  ORData %>%
    filter(city == x) %>%
    ggplot(aes(x = fct_reorder(ethnGroup, logOR), y = logOR, colour = ageGroup)) +
    geom_point(size = 3) +
    geom_errorbar(aes(ymin = lowerCI, ymax = upperCI), width= 0.3) +
    coord_flip() +
    xlab("Ethnic group") +
    ylab("Log OR")
}
```

See Figure 2 in the paper.


Moscow

```
plotOR("Moscow")
```


Rostov

```
plotOR("Rostov")
```




Kazan

```
plotOR("Kazan")
```




Ufa

```
plotOR("Ufa")
```




Vladikavkaz

```
plotOR("Vladikavkaz")
```



Makhachkala

```
plotOR("Makhachkala")
```



Symmetrical odds ratios

Symmetrical odds ratios are odds ratios that involve a pair of ethnic groups. For example, the odds of a Russian woman marrying a Russian rather than a Tatar man divided by the odds of a Tatar woman marrying a Russian rather than a Tatar man. See an application in the social mobility research in Bukodi and Goldthorpe. (2019). Social Mobility and Education in Britain, ch. 4.

See Appendix B in the paper.

For each city, we want to loop over all pairs of ethnic groups and calculate the odds ratio.

A function to produce a data frame of log odds ratios for each pair of ethnicities.

```

symmOR <- function(df){
  # create an empty data frame to store results
  result <- tibble(ethn.wife = character(),
    ethn.husband = character(),
    logOR = numeric())

  k <- 1
  # loop over ethnicities of wives and husbands
  for (i in levels(as.factor(df$ethn.wife))) {
    for (j in levels(as.factor(df$ethn.husband))){
      if (i != j) {
        # calculate and save log OR
        res <- df %>%
          filter(ethn.wife == i | ethn.wife == j) %>%

```

```


        filter(ethn.husband == i | ethn.husband == j) %>%
        collapseData() %>%
        oddsRatio()
      result[k, 1] <- i
      result[k, 2] <- j
      result[k, 3] <- res[1]
      k <- k + 1
    }
  }
}
# remove duplicates (ORs are symmetrical)
result <- result %>%
  mutate(key = paste0(pmin(ethn.wife, ethn.husband),
                        pmax(ethn.wife, ethn.husband), sep = "")) %>%
  distinct(key, .keep_all = TRUE) %>%
  select(-key) %>%
  # re-order factors as in the original data
  mutate(ethn.wife = factor(ethn.wife, levels = levels(as.factor(df$ethn.wife)))) %>%
  mutate(ethn.husband = factor(ethn.husband, levels = levels(as.factor(df$ethn.husband))))
}

# a function to plot the symmetrical odds ratios
plotSymmOR <- function(df){
  df %>%
    symmOR() %>%
    ggplot(aes(x = ethn.wife, y = ethn.husband, fill = logOR)) +
      geom_tile() +
      geom_text(aes(label = round(logOR, 1))) +
      xlab("") +
      ylab("") +
      scale_fill_gradient(low = "white", high = "blue", limits = c(1, 14)) +
      theme_classic() +
      theme(axis.text.x = element_text(angle = 45, hjust = 1),
            axis.line = element_blank(),
            axis.ticks = element_blank())
}

```


Moscow

```
plotSymmOR(Moscow)
```




Rostov

```
plotSymmOR(Rostov)
```


Kazan

```
plotSymmOR(Kazan)
```




Ufa

```
plotSymmOR(Ufa)
```




Vladikavkaz

```
plotSymmOR(Vladikavkaz)
```



Makhachkala

```
plotSymmOR(Makhachkala)
```



Functions to fit the models

For each city I do the following:

- 1) Fit the constant intermarriage rates model (WH + WA + HA).
- 2) Fit the unidiff model (WA + HA + beta*WH).
- 3) Compare G2, BIC, dissimilarity index (on the scale from 0 to 100).

```
# a function to fit three models and produce a table with goodness-of-fit statistics
modelIntermarriage <- function(df){
  set.seed(15)
  # fit the models
  # constant
  constant <- glm(Freq ~ ethn.wife*age.wife + age.wife*ethn.husband + ethn.wife*ethn.husband,
    family = poisson, data = df)

  # unidiff
  unidiff<- glm(Freq ~ ethn.wife*age.wife + ethn.husband*age.wife
    + Mult(Exp(age.wife),
    ethn.wife:ethn.husband), family = poisson, ofInterest = "[.]age.wife",
    data = df, verbose = FALSE)
  # a tibble to store the results, with three rows
  res <- tibble(
    model = c("constant", "unidiff", "unidiff vs. constant"),
    G2 = numeric(3),
    pvalue = numeric(3),
    BIC = numeric(3),
```

```

    diss = numeric(3)
  )
  # populating the tibble
  res$G2 <- c(deviance(constant), deviance(unidiff), lrtest(unidiff, constant)$Chisq[2])
  res$pvalue <- c(NA, NA, lrtest(unidiff, constant)$Pr(>Chisq)[2])
  res$BIC <- round(c(BIC(constant), BIC(unidiff), NA))

  # calculating dissimilarity indices
  dissdif <- df %>%
  mutate(fittedConstant = constant$fitted) %>%
  mutate(diffConstant = abs(fittedConstant - Freq)) %>%
  mutate(fittedUnidiff = unidiff$fitted) %>%
  mutate(diffUnidiff = abs(fittedUnidiff - Freq)) %>%
  summarise(
    dissConstant = sum(diffConstant) / (2 * sum(Freq)) * 100,
    dissUnidiff = sum(diffUnidiff) / (2 * sum(Freq)) * 100
  )

  res$diss <- c(dissdif[[1,1]], dissdif[[1,2]], NA)
  return(res)
}

# a function to estimate unidiff contrasts and return coefficients with quasi standard errors
collectUnidiff <- function(df){
  set.seed(15)
  # re-estimate unidiff
  unidiff <- glm(Freq ~ ethn.wife*age.wife + ethn.husband*age.wife + Mult(Exp(age.wife),
    ethn.wife:ethn.husband), family = poisson,
    ofInterest = "[.]age.wife", data = df, verbose = FALSE)
  # get contrasts
  myContrasts <- getContrasts(unidiff, ofInterest(unidiff))
  unidiffContrasts <- tibble(
    ageGroup = levels(as.factor(df$age.wife)),
    coef = myContrasts$qvframe$estimate,
    se = myContrasts$qvframe$quasiSE,
    city = df$city[[1]]) %>%
    mutate(ageGroup = fct_relevel(ageGroup, ">50", "36-50", "16-35"))
  return(unidiffContrasts)
}

# Unidiff models cannot be reliably estimated for Makhachkala because of the sparseness of the data set
# I recode the data for Makhachkala to a smaller number of groups, keeping 6 largest groups only.

Makhachkala <- Makhachkala %>%
  mutate(ethn.wife = fct_recode(ethn.wife,
    other = "Tabasaran",
    other = "Nogai",
    other = "Rutulian",
    other = "Aghul",
    other = "Azerbaijani",
    other = "Tsakhur")) %>%
  mutate(ethn.husband = fct_recode(ethn.husband,

```

```

        other = "Tabasaran",
        other = "Nogai",
        other = "Rutulian",
        other = "Aghul",
        other = "Azerbaijani",
        other = "Tsakhur")) %>%

group_by(ethn.wife, ethn.husband, age.wife) %>%
summarise(
  Freq = sum(Freq)
) %>%
mutate(city = "Makhachkala") %>%
ungroup()

# Update sixCities
sixCities <- list(Moscow, Rostov, Kazan, Ufa, Vladikavkaz, Makhachkala)
names(sixCities) <- c("Moscow", "Rostov", "Kazan", "Ufa", "Vladikavkaz", "Makhachkala")

```

Models for six cities

Procucing summary statistics for the models and a plot with the unidiff coefficients.

This is Table 2 and Figure 3 in the paper.

For each city, the tables show:

- 1) the constant intermarriage rate model (the intermarriage rates do not change over time),
- 2) the unidiff model (the intermarriage rates change over time),
- 3) the comparison between 2) and 1).

```

sixCities %>%
  map(modelIntermarriage)

## $Moscow
## # A tibble: 3 x 5
##   model          G2    pvalue    BIC    diss
##   <chr>        <dbl>    <dbl> <dbl> <dbl>
## 1 constant      863.    NA      1962  0.190
## 2 unidiff       824.    NA      1933  0.154
## 3 unidiff vs. constant 38.5  4.27e-9    NA  NA
##
## $Rostov
## # A tibble: 3 x 5
##   model          G2    pvalue    BIC    diss
##   <chr>        <dbl>    <dbl> <dbl> <dbl>
## 1 constant      200.    NA      673  0.360
## 2 unidiff       184.    NA      664  0.285
## 3 unidiff vs. constant 16.1  0.000319    NA  NA
##
## $Kazan
## # A tibble: 3 x 5
##   model          G2    pvalue    BIC    diss
##   <chr>        <dbl>    <dbl> <dbl> <dbl>
## 1 constant     3331.    NA    3834  3.01
## 2 unidiff       190.    NA     701  0.300
## 3 unidiff vs. constant 3141.    0     NA  NA

```

```
##
## $Ufa
## # A tibble: 3 x 5
##   model          G2 pvalue    BIC  diss
##   <chr>        <dbl> <dbl> <dbl> <dbl>
## 1 constant    5513.    NA  6921  4.37
## 2 unidiff     1684.    NA  3102  1.93
## 3 unidiff vs. constant 3829.    0    NA  NA
##
## $Vladikavkaz
## # A tibble: 3 x 5
##   model          G2    pvalue    BIC  diss
##   <chr>        <dbl>    <dbl> <dbl> <dbl>
## 1 constant    1100. NA        1818  2.84
## 2 unidiff      254. NA        981  0.839
## 3 unidiff vs. constant 846. 2.04e-184    NA  NA
##
## $Makhachkala
## # A tibble: 3 x 5
##   model          G2    pvalue    BIC  diss
##   <chr>        <dbl>    <dbl> <dbl> <dbl>
## 1 constant    1538. NA        3051  1.67
## 2 unidiff      578. NA        2101  0.640
## 3 unidiff vs. constant 960. 3.56e-209    NA  NA

unidiff.df <- sixCities %>%
  map(collectUnidiff) %>%
  bind_rows()

# plot unidiff contrasts for all six cities
unidiff.df %>%
  ggplot(aes(x = ageGroup, y = coef)) +
  geom_point() +
  geom_errorbar(aes(ymin = coef - 2*se, ymax = coef + 2*se), width = 0.3) +
  geom_hline(yintercept = 0, colour = "red") +
  theme(axis.text.x = element_text(angle = 45, hjust = 1)) +
  ylab("Unidiff coefficient") +
  facet_wrap(~ city) +
  theme(axis.text = element_text(size = 16),
        axis.title = element_text(size = 16, face="bold"))
```